

Apache Spark—Real Time Project—Marketing Analysis from Bank Campaign Data

Solution Code:

The following commands need to be run on the Spark Shell of Scala by given order.

Note:

1. Input data file named “Bank_Data.csv” need to be placed in root HDFS directory.
2. Comment lines like “-- Solution: 1” should be skipped.
3. If required the result can be saved to HDFS directory using given save command `(df2.write.format("com.databricks.spark.csv").save("Result"))`

Commands:

```
spark-shell --master local --packages com.databricks:spark-csv_2.10:1.3.0

case class Bank(age:String,job:String,marital:String,balance:String,y:String)

val rdd = sc.textFile("Bank_Data.csv")

val header = rdd.first()

val data = rdd.filter(row => row != header)

-- Solution: 1 (Creating Data Frame)

val df = data.map(x => x.replaceAll("\"", "").split(";")).map(x=> Bank(x(0),x(1),x(2),x(5),x(16))).toDF()

df.show

df.registerTempTable("temp")

-- Solution: 2 (Subscription Success Rate)

val df2 = sqlContext.sql("select cast(avg(case when t.y == 'yes' then 1.0 else 0 end) as decimal(2,2))
as Success_Rate from temp t")

df2.show
```

-- Solution: 2A (Subscription Failure Rate)

```
val df2 = sqlContext.sql("select cast(avg(case when t.y != 'yes' then 1.0 else 0 end) as decimal(2,2))  
as Failure_Rate from temp t")
```

```
df2.show
```

-- Solution: 3 (Maximum, Average, Minimum Age)

```
val df2 = sqlContext.sql("select max(age) as Maximum_Age, ( ( max(age) + min(age) ) / 2) as  
Mean_Age , round(avg(age)) as Average_Age, min(age) as Minimum_Age from temp")
```

```
df2.show
```

-- Solution: 4 (Average and Median Balance)

```
val df2 = sqlContext.sql("SELECT balance, CAST(AVG(balance) OVER( ) as decimal(16,2)) AS  
Average_Balance, CAST(balance - AVG(balance) OVER ( ) as decimal(16,2) ) AS Difference FROM  
temp")
```

```
df2.show
```

```
df2.write.format("com.databricks.spark.csv").save("Result4")
```

-- Solution: 5 (Subscription by Age)

```
val df2 = sqlContext.sql("select age, count(*) as Subscription_Count from temp where y = 'yes'  
group by age order by age")
```

```
df2.show
```

```
df2.write.format("com.databricks.spark.csv").save("Result5")
```

-- Solution: 6 (Subscription by Marital)

```
val df2 = sqlContext.sql("select marital, count(*) as Subscription_Count from temp where y = 'yes'  
group by marital order by marital")
```

```
df2.show
```

-- Solution: 7 (Subscription by Age & Marital)

```
val df2 = sqlContext.sql("select age, marital, count(*) as Subscription_Count from temp where y =  
'yes' group by age,marital order by age,marital")
```

```
df2.show
```

```
df2.write.format("com.databricks.spark.csv").save("Result7")
```

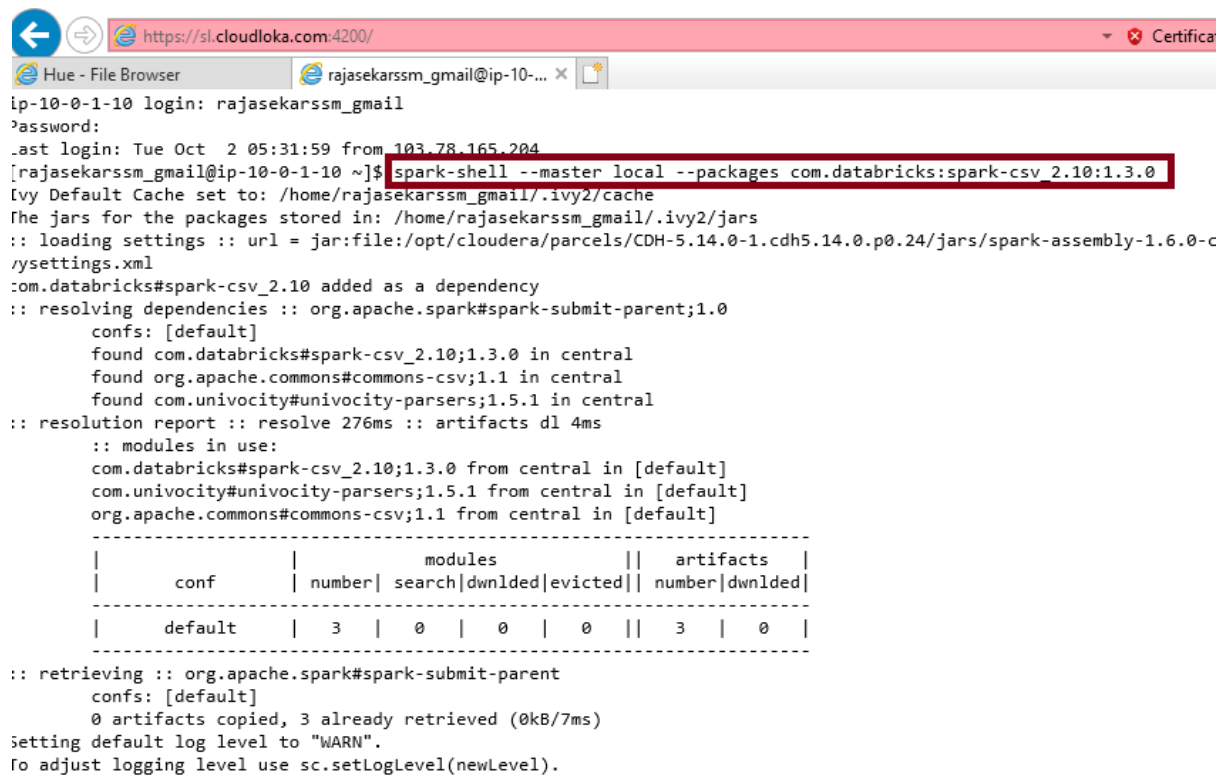
-- Solution: 8 (Feature by Age)

```
val df2 = sqlContext.sql("select sum(case when t.y == 'yes' and age >= 18 and age <= 30 then 1 else 0
end) as Subscriber_Age_18to30, sum(case when t.y == 'yes' and age >= 31 and age <= 60 then 1 else
0 end) as Subscriber_Age_31to60, sum(case when t.y == 'yes' and age > 60 then 1 else 0 end) as
Subscriber_Age_Above60 from temp t")
```

df2.show

Solution Screenshots:

Screenshot 1:



```
ip-10-0-1-10 login: rajasekarssm_gmail
Password:
Last login: Tue Oct  2 05:31:59 from 103.78.165.204
[rajasekarssm_gmail@ip-10-0-1-10 ~]$ spark-shell --master local --packages com.databricks:spark-csv_2.10:1.3.0
Ivy Default Cache set to: /home/rajasekarssm_gmail/.ivy2/cache
The jars for the packages stored in: /home/rajasekarssm_gmail/.ivy2/jars
:: loading settings :: url = jar:file:/opt/cloudera/parcels/CDH-5.14.0-1.cdh5.14.0.p0.24/jars/spark-assembly-1.6.0-c
/ysettings.xml
com.databricks#spark-csv_2.10 added as a dependency
:: resolving dependencies :: org.apache.spark#spark-submit-parent;1.0
  confs: [default]
  found com.databricks#spark-csv_2.10;1.3.0 in central
  found org.apache.commons#commons-csv;1.1 in central
  found com.univocity#univocity-parsers;1.5.1 in central
:: resolution report :: resolve 276ms :: artifacts dl 4ms
  :: modules in use:
  com.databricks#spark-csv_2.10;1.3.0 from central in [default]
  com.univocity#univocity-parsers;1.5.1 from central in [default]
  org.apache.commons#commons-csv;1.1 from central in [default]
-----
|               | modules | artifacts |
| conf | number | search | dwnlded | evicted | number | dwnlded |
-----
| default | 3 | 0 | 0 | 0 | 3 | 0 |
-----
:: retrieving :: org.apache.spark#spark-submit-parent
  confs: [default]
  0 artifacts copied, 3 already retrieved (0kB/7ms)
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel).
```

Screenshot 2:

```
Setting default log level to WARN.
To adjust logging level use sc.setLogLevel(newLevel).
Welcome to

  ____  _
 / ___|| | | |
| |___| |_| |
 \___ \|  _/
      |_|_|

 version 1.6.0

Using Scala version 2.10.5 (Java HotSpot(TM) 64-Bit Server VM, Java 1.8.0_144)
Type in expressions to have them evaluated.
Type :help for more information.
18/10/02 08:32:12 WARN util.Utils: Service 'SparkUI' could not bind on port 40001. Attempting port 40002.
18/10/02 08:32:12 WARN util.Utils: Service 'SparkUI' could not bind on port 40002. Attempting port 40003.
18/10/02 08:32:12 WARN util.Utils: Service 'SparkUI' could not bind on port 40003. Attempting port 40004.
18/10/02 08:32:12 WARN util.Utils: Service 'SparkUI' could not bind on port 40004. Attempting port 40005.
18/10/02 08:32:12 WARN util.Utils: Service 'SparkUI' could not bind on port 40005. Attempting port 40006.
Spark context available as sc (master = local, app id = local-1538469132727).
SQL context available as sqlContext.

scala> case class Bank(age:String,job:String,marital:String,balance:String,y:String)
defined class Bank

scala> val rdd = sc.textFile("Bank_Data.csv")
rdd: org.apache.spark.rdd.RDD[String] = Bank_Data.csv MapPartitionsRDD[1] at textFile at <console>:27

scala> val header = rdd.first()
header: String = "age","job","marital","education","default","balance","housing","loan","outcome","y"

scala> val data = rdd.filter(row => row != header)
data: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[2] at filter at <console>:31

scala>
```

Screenshot 3:

```
scala> val df2 = sqlContext.sql("select cast(avg(case when t.y == 'yes' as Failure_Rate) as decimal(2,2)) as Failure_Rate from Bank t")
df2: org.apache.spark.sql.DataFrame = [Failure_Rate: decimal(2,2)]

scala> df2.show
+-----+
|Failure_Rate|
+-----+
|         0.12|
+-----+

scala> val df2 = sqlContext.sql("select cast(avg(case when t.y != 'yes' as Success_Rate) as decimal(2,2)) as Success_Rate from Bank t")
df2: org.apache.spark.sql.DataFrame = [Success_Rate: decimal(2,2)]

scala> df2.show
+-----+
|Success_Rate|
+-----+
|         0.88|
+-----+
```

Screenshot 4:

```
scala> val df2 = sqlContext.sql("select max(age) as Maximum_Age, ( ( max(age)
df2: org.apache.spark.sql.DataFrame = [Maximum_Age: string, Mean_Age: double,

scala> df2.show
+-----+-----+-----+-----+
|Maximum_Age|Mean_Age|Average_Age|Minimum_Age|
+-----+-----+-----+-----+
|          95|    56.5|        41.0|          18|
+-----+-----+-----+-----+
```

Screenshot 5:

```
scala> val df2 = sqlContext.sql("SELECT balance, CAST(AVG(balance) OVER( ) as decimal(16,2)) AS Average_Balance,
FROM temp")
df2: org.apache.spark.sql.DataFrame = [balance: string, Average_Balance: decimal(16,2), Difference: decimal(16,2)]

scala> df2.show
18/10/02 08:50:48 WARN execution.Window: No Partition Defined for Window operation! Moving all data to window of size 1
+-----+-----+-----+
|balance|Average_Balance|Difference|
+-----+-----+-----+
| 2143|    1362.27|    780.73|
|   29|    1362.27|   -1333.27|
|    2|    1362.27|  -1360.27|
| 1506|    1362.27|    143.73|
|    1|    1362.27|  -1361.27|
|   231|    1362.27|  -1131.27|
|   447|    1362.27|   -915.27|
|    2|    1362.27|  -1360.27|
|   121|    1362.27|  -1241.27|
|   593|    1362.27|   -769.27|
|   270|    1362.27|  -1092.27|
|   390|    1362.27|   -972.27|
|    6|    1362.27|  -1356.27|
|    71|    1362.27|  -1291.27|
|   162|    1362.27|  -1200.27|
|   229|    1362.27|  -1133.27|
|    13|    1362.27|  -1349.27|
|    52|    1362.27|  -1310.27|
|    60|    1362.27|  -1302.27|
|    0|    1362.27|  -1362.27|
+-----+-----+-----+
only showing top 20 rows
```

```
scala> df2.write.format("com.databricks.spark.csv").save("Result4")
18/10/02 08:51:17 WARN execution.Window: No Partition Defined for Window operation! Moving all data to window of size 1
```

Screenshot 6:

```
scala> val df2 = sqlContext.sql("select age, count(*) as Subscription_Count")
df2: org.apache.spark.sql.DataFrame = [age: string, Subscription_Count: long]
```

```
scala> df2.show
```

```
+---+-----+
|age|Subscription_Count|
+---+-----+
| 18|                7|
| 19|               11|
| 20|               15|
| 21|               22|
| 22|              40|
| 23|              44|
| 24|              68|
| 25|             113|
| 26|             134|
| 27|             141|
| 28|             162|
| 29|             171|
| 30|            217|
| 31|            206|
| 32|            221|
| 33|            210|
| 34|            198|
| 35|            209|
| 36|            195|
| 37|            170|
+---+-----+
```

Screenshot 7:

```
scala> val df2 = sqlContext.sql("select marital, count(*) as Subscription_Count")
df2: org.apache.spark.sql.DataFrame = [marital: string, Subscription_Count: long]
```

```
scala> df2.show
```

```
+---+-----+
|marital|Subscription_Count|
+---+-----+
|divorced|            622|
|married|           2755|
|single|           1912|
+---+-----+
```

Screenshot 8:

```
scala> val df2 = sqlContext.sql("select age, marital, c  
df2: org.apache.spark.sql.DataFrame = [age: string, mar:
```

```
scala> df2.show
```

```
+---+-----+-----+  
|age| marital|Subscription_Count|  
+---+-----+-----+  
| 18|  single|          7|  
| 19|  single|         11|  
| 20| married|          1|  
| 20|  single|         14|  
| 21| married|          1|  
| 21|  single|         21|  
| 22|  single|         40|  
| 23| married|          2|  
| 23|  single|         42|  
| 24| married|         10|  
| 24|  single|         58|  
| 25| married|         14|  
| 25|  single|        99|  
| 26| married|         13|  
| 26|  single|       121|  
| 27|divorced|          2|  
| 27| married|         29|  
| 27|  single|       110|  
| 28|divorced|          4|  
| 28| married|         20|  
+---+-----+-----+
```

Screenshot 9:

```
scala> val df2 = sqlContext.sql("select sum(case when t.y == 'yes' and a  
nd age <= 60 then 1 else 0 end) as Subscriber_Age_31to60, sum(case when  
df2: org.apache.spark.sql.DataFrame = [Subscriber_Age_18to30: bigint, Su
```

```
scala> df2.show
```

```
+-----+-----+-----+  
|Subscriber_Age_18to30|Subscriber_Age_31to60|Subscriber_Age_Above60|  
+-----+-----+-----+  
|          1145|          3642|          502|  
+-----+-----+-----+
```